

Patient Health Risk Analytics

Healthcare Analytics Dataset — Question 2
Prepared using Python (pandas, matplotlib, seaborn)

Healthcare Analytics Dataset — Patient Health Risk EDA

Question: A hospital wants to analyze patient health risk based on age, blood sugar, blood pressure, BMI, and disease status. Perform EDA to find important health patterns.

Author: Data Analytics Assignment **Dataset:** patient_health.csv (20 patient records, 6 columns)

Common EDA Tasks Checklist (per updated assessment rubric)

This notebook explicitly covers the following 8 standard EDA tasks:

- 1. **Number of rows and columns** — Task 1 & 2 section
- 2. **Data types of each column** — Task 1 & 2 section
- 3. **Missing values and duplicate values** — Task 3a / 3b / 3c sections
- 4. **Mean, median, minimum, maximum, and standard deviation** — Task 4 section
- 5. **Bar chart, histogram, scatter plot, and box plot** (plus pie chart, line chart, and heatmap as bonus visualizations) — Task 5a-5e sections
- 6. **Important patterns in the dataset** — Task 6 sections (group comparisons and correlation heatmap)
- 7. **4-5 observations based on EDA** — Task 7 section
- 8. **Short conclusion from the analysis** — Task 8 section

Rubric Criteria Alignment (Weightage)

Rubric Criteria		Weightage	Where it's covered
Dataset Understanding	20		Section 2 (dataset description, features, purpose)
Data Cleaning	25		Tasks 3a-3c (missing values, duplicates, inconsistent/incorrect entries)
Data Analysis and Statistical Summary	20		Task 4 (descriptive statistics) + Task 6 (patterns)
Data Visualization	20		Task 5a-5e (bar, pie, histogram, box, scatter, line, heatmap — all labeled)
Insight Generation and	15		Task 7 (observations) + Task 8

1. Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style("whitegrid")
plt.rcParams['figure.figsize'] = (8, 5)
plt.rcParams['axes.titlesize'] = 13
plt.rcParams['axes.titleweight'] = 'bold'
plt.rcParams['axes.labelsize'] = 11

pd.set_option('display.max_columns', None)
print("Libraries imported successfully.")

def show_df(dframe, n=None):
    """Pretty-print a DataFrame (use IPython display() instead in Jupyter)."""
    print((dframe.head(n) if n else dframe).to_string())
```

Libraries imported successfully.

2. Load the Dataset

```
df = pd.read_csv("data/patient_health.csv")
show_df(df, 20)
```

	Patient_ID	Age	Sugar_Level	BP	BMI	Disease_Status
0	P01	45	145	130	28	Yes
1	P02	32	98	118	23	No
2	P03	55	180	145	31	Yes
3	P04	28	90	110	22	No
4	P05	60	200	150	34	Yes
5	P06	40	120	125	26	No
6	P07	52	165	140	30	Yes
7	P08	35	105	120	24	No
8	P09	48	155	135	29	Yes
9	P10	30	95	115	23	No
10	P11	62	210	155	35	Yes
11	P12	38	115	122	25	No
12	P13	50	170	142	30	Yes
13	P14	27	88	108	21	No
14	P15	58	190	148	33	Yes

15	P16	42	130	128	27	No
16	P17	46	150	132	28	Yes
17	P18	34	100	116	24	No
18	P19	53	175	144	31	Yes
19	P20	29	92	112	22	No

Dataset Description

Item	Detail
Records (rows)	20 patients
Features (columns)	6
Patient_ID	Object (string) — unique identifier for each patient
Age	Integer — patient age in years
Sugar_Level	Integer — blood sugar level (mg/dL)
BP	Integer — systolic blood pressure (mmHg)
BMI	Integer — Body Mass Index (kg/m ²)
Disease_Status	Object (categorical) — Yes / No , whether the patient has been diagnosed with the disease being studied

Purpose & significance: This dataset allows the hospital to explore how routinely-collected vitals (age, blood sugar, blood pressure, BMI) relate to a patient’s disease status. Identifying these patterns supports **early risk screening**, helps prioritise patients for preventive care, and gives clinicians a data-driven basis for flagging high-risk profiles before symptoms escalate — a common use case for conditions such as diabetes or hypertension-related illness.

Task 1 & 2: Number of Rows/Columns and Data Types

```
print("Shape (rows, columns):", df.shape)
print()
print("Column data types:")
print(df.dtypes)
```

```
Shape (rows, columns): (20, 6)
```

```
Column data types:
```

```
Patient_ID      str
Age             int64
Sugar_Level     int64
BP              int64
BMI             int64
Disease_Status  str
dtype: object
```

```
print("General info:")
df.info()
```

```
General info:
```

```
<class 'pandas.DataFrame'>
RangeIndex: 20 entries, 0 to 19
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Patient_ID      20 non-null    str
 1   Age             20 non-null    int64
 2   Sugar_Level     20 non-null    int64
 3   BP              20 non-null    int64
 4   BMI             20 non-null    int64
 5   Disease_Status  20 non-null    str
dtypes: int64(4), str(2)
memory usage: 1.1 KB
```

Task 3a: Missing Values

```
missing = df.isnull().sum()
print("Missing values per column:")
print(missing)
print()
print("Total missing values in dataset:", df.isnull().sum().sum())
```

```
Missing values per column:
Patient_ID      0
Age             0
Sugar_Level     0
BP              0
BMI             0
Disease_Status  0
dtype: int64

Total missing values in dataset: 0
```

Observation: The dataset has **no missing values** in any column. No imputation is required here. (In a real hospital dataset we would typically impute missing vitals with the column median rather than drop patient records, since every record matters clinically.)

Task 3b: Duplicate Values

```
dup_rows = df.duplicated().sum()
dup_ids = df['Patient_ID'].duplicated().sum()
print("Fully duplicated rows:", dup_rows)
print("Duplicated Patient_IDs:", dup_ids)

before = df.shape[0]
df = df.drop_duplicates()
after = df.shape[0]
print(f"Rows before: {before}, after de-duplication: {after}")
```

```
Fully duplicated rows: 0
Duplicated Patient_IDs: 0
Rows before: 20, after de-duplication: 20
```

Observation: No duplicate rows or duplicate Patient IDs were found; each patient appears exactly once.

Task 3c: Data Cleaning — Consistency Checks (Incorrect Entries & Range Validation)

```
# Standardise text fields
df['Patient_ID'] = df['Patient_ID'].str.strip()
df['Disease_Status'] = df['Disease_Status'].str.strip().str.capitalize()

print("Unique values in 'Disease_Status':", df['Disease_Status'].unique().tolist())

# Clinical range/sanity checks on numeric columns
checks = {
    'Age (0-120)': df['Age'].between(0, 120).all(),
    'Sugar_Level (40-400 mg/dL)': df['Sugar_Level'].between(40, 400).all(),
    'BP (70-220 mmHg)': df['BP'].between(70, 220).all(),
    'BMI (10-60 kg/m2)': df['BMI'].between(10, 60).all(),
}
for label, ok in checks.items():
    print(f"{label}: {'OK' if ok else 'OUT OF RANGE VALUES FOUND'}")

# Convert Disease_Status to a proper categorical dtype
df['Disease_Status'] = df['Disease_Status'].astype('category')
print()
print("Disease_Status column dtype now:", df['Disease_Status'].dtype)
```

```
Unique values in 'Disease_Status': ['Yes', 'No']
Age (0-120): OK
Sugar_Level (40-400 mg/dL): OK
BP (70-220 mmHg): OK
BMI (10-60 kg/m2): OK

Disease_Status column dtype now: category
```

Observation: All vitals fall within clinically plausible ranges — no negative ages, implausible sugar levels, or impossible BP/BMI values. `Disease_Status` contains exactly two clean categories (Yes/No) and has been converted to a `category` dtype.

Task 4: Descriptive Statistics — Mean, Median, Min, Max, Std

```
num_cols = ['Age', 'Sugar_Level', 'BP', 'BMI']
print("Descriptive statistics for numerical columns:")
show_df(df[num_cols].describe().round(2))
```

Descriptive statistics for numerical columns:

	Age	Sugar_Level	BP	BMI
count	20.00	20.00	20.00	20.00

mean	43.20	138.65	129.75	27.30
std	11.32	40.33	14.52	4.26
min	27.00	88.00	108.00	21.00
25%	33.50	99.50	117.50	23.75
50%	43.50	137.50	129.00	27.50
75%	52.25	171.25	142.50	30.25
max	62.00	210.00	155.00	35.00

```
print("Additional statistics:")
summary = pd.DataFrame({
    'mean': df[num_cols].mean(),
    'median': df[num_cols].median(),
    'std': df[num_cols].std(),
    'skew': df[num_cols].skew(),
})
show_df(summary.round(2))
```

Additional statistics:

	mean	median	std	skew
Age	43.20	43.5	11.32	0.10
Sugar_Level	138.65	137.5	40.33	0.26
BP	129.75	129.0	14.52	0.14
BMI	27.30	27.5	4.26	0.22

Task 5a: Bar Chart & Pie Chart — Categorical Analysis (Disease_Status)

```
status_counts = df['Disease_Status'].value_counts()
status_pct = df['Disease_Status'].value_counts(normalize=True).mul(100).round(1)

print("Counts:")
print(status_counts)
print()
print("Percentage:")
print(status_pct.astype(str) + '%')
```


```
Counts:
Disease_Status
No      10
Yes      10
Name: count, dtype: int64

Percentage:
Disease_Status
No      50.0%
Yes      50.0%
Name: proportion, dtype: str
```

```

colors = ['#D64550', '#4C956C'] # Yes = red (risk), No = green (healthy)
fig, ax = plt.subplots()
bars = ax.bar(status_counts.index, status_counts.values, color=colors, edgecolor='black')
for b in bars:
    ax.text(b.get_x()+b.get_width()/2, b.get_height()+0.15, str(int(b.get_height())) ,
            ha='center', va='bottom', fontweight='bold')
ax.set_title('Bar Chart: Number of Patients by Disease Status')
ax.set_xlabel('Disease Status')
ax.set_ylabel('Number of Patients')
plt.tight_layout()
plt.show()

```




```

fig, ax = plt.subplots()
ax.pie(status_counts.values, labels=status_counts.index, autopct='%1.1f%%',
       colors=colors, startangle=90, wedgeprops={'edgecolor':'white'})
ax.set_title('Pie Chart: Proportion of Patients With vs Without Disease')
ax.legend(status_counts.index, title="Disease Status", loc="best")
plt.tight_layout()
plt.show()

```

Pie Chart: Proportion of Patients With vs Without Disease




Observation: 10 of 20 patients (50%) are diagnosed positive (Yes) and 10 (50%) are negative (No) — a perfectly balanced sample for comparison.

Task 5b: Histogram & Box Plot — Numerical Distributions

```
fig, axes = plt.subplots(1, 4, figsize=(19, 4.5))
colors_hist = ['#4C72B0', '#DD8452', '#55A868', '#8172B2']

for ax, col, c in zip(axes, num_cols, colors_hist):
    ax.hist(df[col], bins=6, color=c, edgecolor='black', alpha=0.85)
    ax.set_title(f'Histogram: {col} Distribution')
    ax.set_xlabel(col)
    ax.set_ylabel('Frequency')

plt.tight_layout()
plt.show()
```




```
fig, axes = plt.subplots(1, 4, figsize=(19, 4.5))
for ax, col, c in zip(axes, num_cols, colors_hist):
```



```
ax.boxplot(df[col], vert=True, patch_artist=True,
           boxprops=dict(facecolor=c, alpha=0.6))
ax.set_title(f'Box Plot: {col}')
ax.set_ylabel(col)
ax.set_xticks([])
```

```
plt.tight_layout()
plt.show()
```




Observation: No extreme outliers appear across Age, Sugar_Level, BP or BMI — all four vitals show fairly symmetric, moderate spreads consistent with a general adult patient population.

Task 6a: Important Patterns — Vitals vs Disease Status

```
fig, axes = plt.subplots(1, 4, figsize=(19, 4.5))
for ax, col, c in zip(axes, num_cols, colors_hist):
    data_yes = df[df['Disease_Status'] == 'Yes'][col]
    data_no = df[df['Disease_Status'] == 'No'][col]
    bp = ax.boxplot([data_yes, data_no], tick_labels=['Yes', 'No'], patch_artist=True)
    for patch, col2 in zip(bp['boxes'], colors):
        patch.set_facecolor(col2)
        patch.set_alpha(0.6)
    ax.set_title(f'{col} by Disease Status')
    ax.set_ylabel(col)

plt.tight_layout()
plt.show()
```




```
group_means = df.groupby('Disease_Status', observed=True)[num_cols].mean().round(2)
show_df(group_means)
```

	Age	Sugar_Level	BP	BMI
Disease_Status				
Yes				
No				

No	33.5	103.3	117.4	23.7
Yes	52.9	174.0	142.1	30.9

```
fig, ax = plt.subplots(figsize=(9,5))
x = np.arange(len(num_cols))
width = 0.35
yes_vals = group_means.loc['Yes']
no_vals = group_means.loc['No']

ax.bar(x - width/2, yes_vals, width, label='Yes (Disease)', color='#D64550', edgecolor='black')
ax.bar(x + width/2, no_vals, width, label='No (Healthy)', color='#4C956C', edgecolor='black')
ax.set_xticks(x)
ax.set_xticklabels(num_cols)
ax.set_ylabel('Average Value')
ax.set_title('Bar Chart: Average Age / Sugar Level / BP / BMI by Disease Status')
ax.legend(title='Disease Status')
plt.tight_layout()
plt.show()
```



Observation: Patients diagnosed with the disease (Yes) show clearly higher average Sugar_Level, BP and BMI, and are somewhat older on average, compared to patients without the disease (No) — all four vitals trend in the same direction as disease risk.

Task 5c: Scatter Plot — Bivariate Relationships

```
fig, axes = plt.subplots(1, 3, figsize=(16, 5))
pairs = [('Age', 'Sugar_Level'), ('Sugar_Level', 'BP'), ('BMI', 'BP')]


for ax, (xcol, ycol) in zip(axes, pairs):
    for status, c in zip(['Yes', 'No'], colors):
        subset = df[df['Disease_Status'] == status]
        ax.scatter(subset[xcol], subset[ycol], label=status, color=c,
                    edgecolor='black', s=70)
    ax.set_title(f'{xcol} vs {ycol}')
```

```

ax.set_xlabel(xcol)
ax.set_ylabel(ycol)
ax.legend(title='Disease Status')

plt.tight_layout()
plt.show()

```



Observation: Diseased patients (red) cluster toward the higher end of sugar level, blood pressure and BMI in every pairing, while healthy patients (green) cluster toward the lower end — a visibly separable pattern even with simple two-variable views.


Task 5d: Line Chart — Trend Analysis (Bonus Visualization)

```

trend_df = df.sort_values('Age').reset_index(drop=True)

fig, ax1 = plt.subplots(figsize=(11,5))
ax1.plot(trend_df.index, trend_df['Sugar_Level'], marker='o', color='#DD8452', label='Sugar Level')
ax1.plot(trend_df.index, trend_df['BP'], marker='s', color='#4C72B0', label='Blood Pressure (mmHg)')
ax1.plot(trend_df.index, trend_df['BMI']*4, marker='^', color='#55A868', label='BMI (x4 for scale)')
ax1.set_xlabel('Patients sorted by ascending Age')
ax1.set_ylabel('Value')
ax1.set_title('Line Chart: Sugar Level, BP & BMI Trend Across Patients (by Age)')
ax1.legend(loc='upper left')
plt.tight_layout()
plt.show()

```




Observation: As patients are ordered by increasing age, sugar level, blood pressure and BMI all trend broadly upward together — suggesting these risk factors accumulate jointly with age in this sample.

Task 6b: Important Patterns — Correlation Heatmap

```
df_corr = df.copy()
df_corr['Disease_Num'] = (df_corr['Disease_Status'] == 'Yes').astype(int)
corr_matrix = df_corr[num_cols + ['Disease_Num']].corr().round(2)
show_df(corr_matrix)
```

	Age	Sugar_Level	BP	BMI	Disease_Num
Age	1.00	0.99	0.99	1.00	0.88
Sugar_Level	0.99	1.00	0.99	0.99	0.90
BP	0.99	0.99	1.00	0.99	0.87
BMI	1.00	0.99	0.99	1.00	0.87
Disease_Num	0.88	0.90	0.87	0.87	1.00

```
fig, ax = plt.subplots(figsize=(7,6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1,
            square=True, linewidths=0.5, cbar_kws={'label': 'Correlation coefficient'}, ax=ax)
ax.set_title('Heatmap: Correlation Between Vitals and Disease Status')
plt.tight_layout()
plt.show()
```



Observation: `Sugar_Level`, `BP` and `BMI` are all strongly positively correlated with each other and with `Disease_Num` (disease presence). `Age` also correlates positively with disease, but somewhat more weakly than the three direct clinical measurements — consistent with sugar level, blood pressure and `BMI` being the more proximate risk indicators.

Task 7: Observations (4-5 Key Observations)

1. **Dataset is clean:** no missing values, no duplicate records, and all vitals fall within clinically plausible ranges.
2. **Class balance:** exactly 50% of patients are disease-positive and 50% are disease-negative — an ideal balance for comparison.
3. **`Sugar_Level`, `BP` and `BMI` are all strongly positively correlated** with each other and with `Disease_Status`.
4. Patients diagnosed with the disease show markedly higher average sugar level, blood pressure, and `BMI` — scatter plots show largely separable clusters between diseased and healthy patients.
5. **`Age` correlates positively but more moderately** with disease status than the three direct clinical measurements.

Task 8: Conclusion

Sugar level, blood pressure, and `BMI` together form a strong composite risk signal for disease status in this dataset, while age plays a secondary, amplifying role. The hospital can use a combined score across these vitals — rather than any single measurement — to triage patients for

closer monitoring or preventive care. Clinical decision thresholds would need validation on a much larger, representative patient population before real-world use.